

**CIS 4930 Introduction to Hadoop and Big Data**  
**Lab Homework 7**  
**Julie Reyes**

## **1 InvertedIndexInput**

### **1.1 Project Introduction**

The purpose of this project is to perform an inverted index job. Specifically, the job for this project is performed on the `invertIndexInput.tgz` file, which is decompressed and submitted to the Hadoop distributed file system. This project will output an index list for all the words in the text. For example, for an instance of a word found in the text, the index will contain a list of all the locations the word is found.

### **1.2 Driver Description**

The driver for this project is named `InvertedIndex`, and is configured to implement `Tool`. The driver defines the input and output paths to be defined by the command line arguments. The driver map class is set to be the `IndexMapper.class` and the Reducer is set to `IndexReducer.class`. The driver sets the map output key and value class to be of type `Text`. The input format class for this project is set to `KeyValueTextInputFormat.class`. The driver sets the job name as “Inverted Index”. The driver sets the driver class to be `InvertedIndex.class`.

### **1.3 Description on Mapper**

The Mapper for this project is named `IndexMapper.java`. This Mapper extends the `Mapper` class. This Mapper takes in four parameters, all of which are of type `Text`. The `Map` function overrides the `Mapper` and takes in three parameters including the `Text` key, `Text` value, and `Context` context. The `map` function utilizes `FileSplit` object which is initialized by the `context.getInputSplit()` that is cast to a `FileSplit` object. Next, the file path is extracted by the function `getPath()`. The path of the file is saved in a variable named `fileName`, which is modified to include an “@” and the key as the whole string. Therefore, the final `fileName` string is represented as `filename@key`. Before appending the key to the `fileName` string it is converted to type string with the `toString()` function. Next, the variable `line` is initialized as the value converted to type string with `toString()` method and made lowercase with the `toLowerCase()` method. A for loop is used to iterate through the input line and split it into words, which are then emitted by the `context.write()` method. The key emitted from the Mapper is the word extracted from the input line and the value emitted is the filename and the line number the word is found.

### **1.4 Description on Reducer**

The Reducer class for this project is called `IndexReducer`, which extends the `Reducer` class that takes in four parameters of which are all of type `Text`. The `reduce` method overrides the `Reducer`, and takes in three parameters. The parameters the `reduce` method accepts includes, the `Text` key, `Iterable<Text>` values, and `Context` context. The `reduce` method then uses the `StringBuilder` class to create a mutable string named `valueList`. The `valueList` string contains an index of all the locations of the key word. A for loop takes in the iterable values and appends the values to the `valueList` string with a comma separating the filenames. The final output that is emitted from the Reducer is the word received from the Mapper as the key and the value is the index list that includes all the locations of the word by file name and line number it is found at separated by commas.

### **1.5 Data flow description**

Data flow description. The file that the job for this project is performed on is named `invitedIndexInput`. This file is decompressed and placed on the Hadoop distributed file system. The driver configures the job then submits it to the Hadoop Cluster. Once the job is submitted to the Cluster, the tasks are submitted to the data nodes. The driver code determines the format that the data is sent to the Mappers. In this project the input format type was of the type `KeyValueTextInputFormat.class`, which extends the `TextInputFormat`. Therefore, the record reader splits the data into input splits of type `FileSplit`. Next, the data goes to the Mapper from the record reader. In the Mapper the data is broken down by input splits into a keyword value pair that is emitted to the Reducer. In the Reducer the data is modified to emit the keyword from the file and all the index values associated with the word in the form of a list. The final output of the Reducer is saved to the output file specified in the command line.

## **1.6 Test data and results**

Test data and results: The job was tested in the eclipse environment with the Run Configuration option. The test run was saved locally and verified to ensure proper output. Once the job output was correct on eclipse test run, the job was submitted to the Hadoop Cluster. The following is an example of the final output:

```
accessible    cymbeline@2332,  
accidence    merrywivesofwindsor@3016,
```

## **2 Word Co-Concurrency**

### **2.1 Project Introduction**

The purpose of this project is to calculate the number of times words appear next to each other. This may be of importance when determining connections, especially in the field of data mining. For this project the shakespeare file on the Hadoop Cluster from previous labs was used to verify the efficacy of the MapReduce job for this project.

### **2.2 Description on driver**

The driver for the job in this project is named `WordCo.java`. In this driver the class `WordCo` extends `Configured` and implements `Tool`. First, in the driver program, the command line arguments are checked to determine if the proper command has been input. Next, the job is initialized by the `Job(getConfig())` method. The jar for this project is set by the `.setJarByClass()` method as `WordCo.class`. The job name is specified as "Word Co-Occurrence". The file path for the input is specified by the command line argument at `args[0]`, and the file output path is specified by the command line argument at `args[1]`. The Mapper class is set to `WordCoMapper.class` and the Reducer is set to `SumReducer.class`. The output key for this project is of type `Text` and the output value class type is set to `IntWritable.class`. Lastly, the driver sets the job to wait until complete.

### **2.3 Description on Mapper**

The Mapper for this project is defined as `WordCoMapper`. The Mapper class is extended by `WordCoMapper`. The Mapper takes in four parameters of type `LongWritable`, `Text`, `Text`, `Text`, and `IntWritable`. The subclass method `map()` overrides the Mapper function. The `map()` method takes in three parameters that include the `LongWritable` key, `Text` value, and `Context` context. The input value of type `Text` is converted to a lower case string with the methods `toString()` and `toLowerCase()`. Next, a temporary `String` variable `temp` is initialized to null. Since, Mapper is looking for pairs of words that are next to each other the boolean variable `firstWords` is initialized to true. This variable is needed to ensure that the first word is

selected before finding the word that is next to it. Next, a for loop iterates through the line input and breaks it up by pairs words. If the boolean firstWord is false then the String variable temp is initialized as the first word in the pair. The for loop iterates through values to find all the words that are next to each other. The pairs of words are emitted by the context variable with the write() method as the key output in the format (word one, word two) of type Text. The value emitted from the Mapper is of type IntWritable and is just the number one as each unique pair will be counted in the Reducer.

## **2.4 Description on Reducer**

The Reducer for this project is named SumReducer. This Reducer is the same Reducer used in the previous Labs. The SumReducer extends the Reducer class and accepts four parameters including Text, IntWritable, Text, and IntWritable. The reduce() method overrides and takes in the parameters Text key, Iterable<IntWritable> values, and Context context. Next, a variable named wordCount is initialized to zero. This variable is used to keep count of the pairs that were the same. To count all similar pairs a for loop iterates through the values and uses the get() method to access them. The Reducer then emits the final key value pairs, which includes the unique pair of consecutive words and the value is the number of times that pair was found on the file.

## **2.5 Data flow description**

Initially the data for this project is found on the Hadoop Cluster. When the job is submitted the job configured on the local machine. Next, the job is submitted to the Cluster. Once the job is submitted, the data input splits are sent to the Mapper. The Mapper takes the data provided by the record reader and forms key value pairs. For this project the pairs are the keys and the value is just one. The intermediate data is shuffled and sorted. Once, last Mapper is complete and the shuffle and sort phase is complete the Reducers begin the reduce tasks on the intermediate data. Lastly, once the all similar pairs of co-current words are counted, the reduce method outputs the pair as the key and the total number of times it was encountered as the value. The final data output is maintained on the Hadoop Cluster.

## **2.6 Test data and results**

The final job for this project was tested within eclipse and on the Hadoop Cluster. The following is an expert of the final output from the job submitted.

|             |   |
|-------------|---|
| youth,hath  | 2 |
| youth,have  | 1 |
| youth,he    | 5 |
| youth,henry | 1 |
| youth,here  | 1 |
| youth,how   | 1 |